


TREC: do more hypotheses substitute for encoder capacity? (pool trec_gen256, train 20/class, 5 seeds)

accuracy vs pool size (per encoder size)



calibration vs pool size

